

1) Explain ACID properties of a transaction with suitable example.

A database transaction must satisfy four key properties — ACID:

1 Atomicity

A transaction is “all or nothing”.

If any step fails, the whole transaction is rolled back.

Example:

Money transfer from A to B

- Debit A
 - Credit B
If debit succeeds but credit fails → rollback debit.
-

2 Consistency

Transaction must convert database from one valid state to another valid state.

Example:

Balance of A + B must remain same before and after transfer.

3 Isolation

Concurrent transactions should not interfere with each other.

Example:

While T1 is updating A,
T2 should not read intermediate value.

4 Durability

Once committed, transaction results must persist even after system failure.

Example:

If system crashes after commit → balances remain updated.

2) Define Transaction. Explain various states of a transaction.

A **transaction** is a sequence of database operations forming a single logical unit of work.

Example operations:

Read, Write, Update, Insert, Delete

Transaction States

1 Active

Transaction execution in progress.

2 Partially Committed

Final statement executed.

3 Committed

Transaction completed successfully.

4 Failed

Error occurred during execution.

5 Aborted

Transaction rolled back and database restored.

6 Terminated / End

Transaction finished after commit or abort.

3) Explain Two-Phase Commit Protocol (2PC).

Used in **distributed databases** to ensure atomicity across multiple sites.

Phase-1: Prepare / Voting Phase

Coordinator sends **Prepare(T)** to all participants.

Each site replies:

- **Vote-Commit** → ready to commit

- **Vote-Abort** → cannot commit
-

Phase-2: Commit Phase

If *all* vote-commit → Coordinator sends **Global-Commit**

Else → Coordinator sends **Global-Abort**

Failure Handling

- If coordinator fails → participants wait/timeout
 - Recovery is done using logs
-

2PC ensures:

Atomicity
Consistency
Slow due to blocking behaviour

4) Differentiate between Conflict Serializability & View Serializability.

Basis	Conflict Serializability	View Serializability
Definition	Schedules obtained by swapping non-conflicting operations	Schedules that produce same read & write results
Condition	Conflicting operations must preserve order	Final read, write & initial read must match
Test method	Precedence Graph	Difficult to test
Guarantee	Conflict serializable \Rightarrow View serializable	View serializable $\not\Rightarrow$ Conflict serializable

Practical usage	Used in concurrency control	Theoretical concept
-----------------	-----------------------------	---------------------

5) Explain Deadlock with suitable example.

Deadlock occurs when:

Two or more transactions wait for each other's locked resources and none can proceed.

Example

T1 locks A → needs B
T2 locks B → needs A

Result → **circular wait**

Deadlock Conditions

1. Mutual exclusion
 2. Hold and wait
 3. No pre-emption
 4. Circular wait
-

Deadlock Handling

- Deadlock prevention
- Deadlock avoidance (Wait-Die / Wound-Wait)
- Deadlock detection & recovery

6) Write short note on Log-Based Recovery.

DBMS maintains a **log file** to support recovery.

Log Record Format

<TransactionID, DataItem, Old Value, New Value>

Types of Logging

1 Undo Logging

Restores old values during rollback.

2 Redo Logging

Reapplies committed updates after crash.

3 Undo-Redo Logging

Supports both rollback & recovery.

Commit Logging Rules

- <T start> — transaction begins
 - <T, X, V1, V2> — before & after values logged
 - <T commit> — commit completed
-
-

7) Explain Query Processing Steps (or with diagram).

Query processing converts **high-level SQL** into **efficient execution plan**.

Steps

1 Query Parsing & Syntax Checking

- verifies query syntax
- produces parse tree

2 Query Optimization

- chooses lowest-cost execution plan
- uses statistics

3 Execution Plan Generation

4 Query Execution Engine

- executes operations
 - interacts with storage
-
-

8) Explain Heuristics in Query Optimization.

Heuristics are **rule-based transformations** to reduce query cost.

Common Heuristic Rules

1. Perform **selection early**
2. Perform **projection early**
3. Replace Cartesian product with **joins**
4. Perform operations on **smaller relations first**
5. Combine selections & projections

Example

Before optimization

$\sigma \text{ city}=\text{"PUNE"} (\text{Customer} \times \text{Account})$

After heuristic optimization

$\sigma \text{ city}=\text{"PUNE"} (\text{Customer} \bowtie \text{Account})$

Result → fewer tuples processed